Automated Reinforcement Learning: An Overview

Reza Refaei Afshar[a,*], Yingqian Zhang[a], Joaquin Vanschoren[b], Uzay Kaymak[a]

[a]*School of Industrial Engineering, Eindhoven University of Technology, The Netherlands*
[b]*School of Mathematics and Computer Science, Eindhoven University of Technology, The Netherlands*

Abstract

Reinforcement Learning and recently Deep Reinforcement Learning are popular methods for solving sequential decision making problems modeled as Markov Decision Processes. RL modeling of a problem and selecting algorithms and hyper-parameters require careful considerations as different configurations may entail completely different performances. These considerations are mainly the task of RL experts; however, RL is progressively becoming popular in other fields where the researchers and system designers are not RL experts. Besides, many modeling decisions, such as defining state and action space, size of batches and frequency of batch updating, and number of timesteps are typically made manually. For these reasons, automating different components of RL framework is of great importance and it has attracted much attentions in recent years. Automated RL provides a framework in which different components of RL including MDP modeling, algorithm selection and hyper-parameter optimization are modeled and defined automatically. In this article, we explore the literature and present recent work that can be used in automated RL. Moreover, we discuss the challenges, open questions and research directions in AutoRL.

Keywords: `Reinforcement Learning`, `Automated Reinforcement Learning Pipeline`, `AutoRL`, `Sequential Decision Making`

---

*Corresponding author.
Email addresses:* `r.refaei.afshar@tue.nl` (Reza Refaei Afshar), `YQZhang@tue.nl` (Yingqian Zhang), `j.vanschoren@tue.nl` (Joaquin Vanschoren), `U.Kaymak@ieee.org` (Uzay Kaymak)

1. Introduction

Reinforcement Learning (RL) is a learning approach in which no prior knowledge of an environment is necessary. An agent learns the optimal behavior known as *policy* by interacting with the environment. At each decision step, the agent observes the *state* of the environment, takes an *action* and receives a scalar *reward* value from the environment. Using this reward value, the agent adjusts its policy in order to maximize the long term reward. The long term reward is either the sum of all future rewards or the discounted sum of future rewards in order to reduce the impact of future actions [84].

RL is a method to solve sequential decision making problems modeled as Markov Decision Process (MDP). Formally speaking, a discrete time finite horizon MDP is a tuple (S, A, r, T, γ), where S is the state space, A is the action space, $r : S \times A \to r \in \mathbb{R}$ is an instant reward value denoting the benefit of transition from current state $s_t \in S$ to the next state $s_{t+1} \in S$, and γ is the discount factor [59]. At each decision moment or discrete timestep t, an agent interacts with the environment and its goal is to learn a policy $\pi : S \times A \to \pi \in [0, 1]$ that determines a probability value for each action. Following a greedy, ϵ-greedy, softmax or other action selection policies, the agent takes an action according to the probabilities and transitions to the next state. In other words, the agent observes a state $s_t \in S$ and performs an action $a_t \sim \pi(.|s_t)$. Taking an action has two consequences. First the agent receives a reward value r_t. Then, the state of the environment transitions from s_t to a new state s_{t+1} based on the transition probability T. The agent updates the policy during the learning phase in order to find the optimal policy that yields the maximum total reward. RL provides an interaction-based framework to solve the MDP and learn the policy π.

MDPs can be continuous time, infinite horizon or partially observable. Infinite horizon MDPs do not have a goal state and the system runs for ever [84]. The agent learns to maximize its total reward without expecting a goal state. In continuous time MDPs, unlike discrete time, the decisions are made

2

at every point in time. The formulation of continuous time MDP is similar to discrete MDP, however, these problems are harder to solve in general [33]. Partially Observable MDP (POMDP) is a kind of MDP where the agent has limited information about the environment due to different reasons such as limited sensors or uncertainty in the environment [70, 84]. These variants of MDPs could also be modeled and solved by RL, although they need different considerations on defining the states which the agent has a *belief* about the environment rather than complete observation. In this article, we focus on discrete time fully observable MDPs.

To model a problem in RL framework and solve it accordingly, different components of RL, including MDP modeling, algorithms, and hyper-parameters such as the number of training steps, the structure of policy function, etc., should be determined before starting the learning procedure. In common practice, the exact observation from the environment is used as a state representation and the agent's decision is directly sent to the environment as an action. Although these states and actions might help the agent to find an optimal or near optimal policy for some tasks, they are not necessarily the best representations for the states and the actions. In some other tasks, further processing of the environment is required to find a suitable state representation. As a simple example, normalization is usually necessary for the inputs of Neural Networks (NNs) and using raw observation of the environment may produce undesirable results. Furthermore, the performance of a RL policy is highly dependent on the RL algorithms and its hyper-parameters. Selecting algorithms and tuning hyper-parameters are normally done by using expert knowledge. However, many iterations are still needed to find the best set of hyper-parameters. In addition, there are many optimization problems in different areas with no available expert knowledge on RL, that could be solved by RL. Hence, automating an RL procedure in order to facilitate using this approach for non-experts is of great importance.

Automated Reinforcement Learning (AutoRL) provides a framework to automatically make appropriate decisions about the settings of an RL procedure before starting the learning. In other words, RL components, including state,

3

action and reward, algorithm selection, and hyper-parameters optimization, are determined through AutoRL, and the best configuration for each component is provided for an RL procedure to solve a task. Figure 1 shows an RL pipeline containing the RL components. To model and solve a problem using RL, the steps of this pipeline are followed, starting from MDP modeling. AutoRL aims to automate different steps of this pipeline and reduce the necessity of expert knowledge. We use AutoRL to emphasize the resemblance with Automated Machine Learning (AutoML) as a framework for automating supervised and unsupervised learning procedures. According to [35], an AutoML pipeline consists of the components of Machine Learning frameworks. It starts with data preprocessing that contains data collection and data cleaning. The next step is feature engineering including feature extraction, feature construction and feature selection. After processing the data and defining the features, a model is developed and optimized to perform classification or clustering tasks. Finally, the quality of the model is determined through evaluation. In order to automate this procedure, the output of evaluation is used to configure new settings and follow the steps again. In this way, several iterations might be needed to derive the optimal settings for data preprocessing, feature engineering and model generation. Although some components of AutoML and AutoRL are similar, AutoML methods cannot necessarily be used in AutoRL because the configuration of the problems and the complexity of the evaluation step are different. AutoRL pipeline builds a pipeline for RL to define and select components of RL automatically. In recent years, by combining Deep Learning and RL, the need for automating the components of Deep Reinforcement Learning (DRL) is increased because many modeling decisions are made manually and even RL experts have to test several different configurations related to state, action, reward, algorithm and hyper-parameters, to obtain the best definition.

In this article, we review relevant work that can be included in an AutoRL framework and elaborate research challenges and directions in this relatively new research area. For each of the components mentioned above, different approaches in the literature are presented that might be helpful in automating
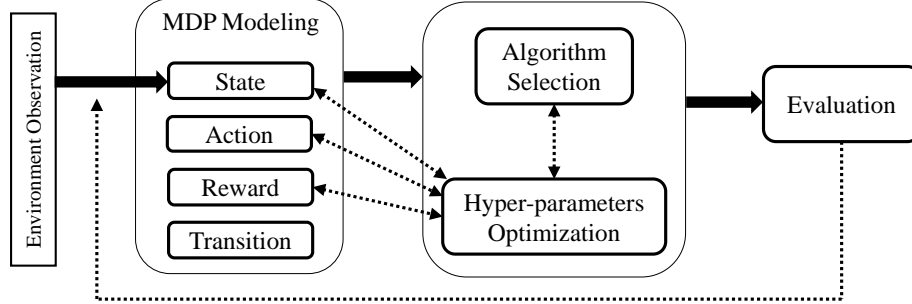
4

Figure 1: Solid arrow shows RL pipeline and dashed arrows depict AutoRL pipeline where the evaluation results are used to set new configurations. This loop continues until finding an appropriate configuration.

the corresponding component. For example, approaches that modify the initial observations of the environment to define a state representation are candidate methods for using in automating states. AutoRL pipeline consists of modeling a particular problem as a sequential decision making problem and MDP, selecting algorithm and tuning hyper-parameters. These three steps are illustrated in Figure 1 followed by evaluation. Evaluating a RL algorithm is normally performed by tracing reward alteration during training and comparing the final total reward with some baselines.

In order to clarify the purpose of AutoRL, assume Traveling Salesman Problems (TSP) as an example. Recently, DRL has been popular for learning heuristics and solving different variants of TSP. The first step in solving TSP with DRL is to model the problem as a sequential decision making and determine MDP components. Vanilla TSP is defined as finding a tour with minimum length in a graph where all the nodes are visited exactly once. In a constructive solution approach, an agent starts from the source and walks through nodes until building a tour. At each timestep, the graph, current node, visited nodes and remained nodes are the observations of an environment. There are several ways like graph neural networks [64] and structure2vec [17] to convert the ob-

servations to state representation and AutoRL helps to find the best approach among the possible methods for this conversion. The second step is to define a RL algorithm for updating the policy network. Popular RL algorithms such as A2C [52], PPO [65], ACER [81] and DQN [53] are available to train the policy network. Each of them is useful for particular problems and it is not easy to determine the best algorithm for a problem. AutoRL aims to provide the opportunity of searching among these possible algorithms to find the suitable one. We focus on model-free RL algorithms where the transition probabilities between states are unknown. The third step is to set the hyper-parameters such as network architecture, learning rate, discount factor, etc. AutoRL employs hyper-parameter optimization methods in its framework to derive the optimal hyper-parameters for an algorithm and a set of problem instances.

This paper is organized as follows. In section 2, previous work related to automating the MDP modeling including the definition of states, actions and rewards are explained. In section 3, the process of RL algorithm selection is reviewed. Since algorithm selection is normally intertwined with hyper-parameter optimization, most of the combined work in algorithm selection and hyper-parameter optimization together with different hyper-parameter optimization work are presented in section 4. Section 5 present recent work in meta-learning that can be leveraged in a RL framework. In section 6, previous work in optimizing and learning neural network architecture are reviewed. Finally, section 7 concludes the paper.

2. Markov Decision Process Components

Since Reinforcement Learning aims to solve MDPs, an AutoRL agent needs to define these four components: state, action, reward, transition probability. The transition probability is mostly unknown for model-free RL problems [74]. Hence, we focus only on state, action and reward definition in this section.

2.1. Methods for Automating States

As mentioned before, the original observation that an agent receives from interacting with environment is commonly assumed as a state. This holds for many classical RL problems such as mountain car, cart pole and pendulum [84]. Nevertheless, the original observation is not necessarily an efficient state representation. Moreover, in some applications, especially for those with continuous state space, the state space is extremely large and approximating a good value functions is difficult. Hence, deriving a particular mapping from observations to states that helps the final policy take better actions, attracts much attention in the literature. We categorize the existing work in defining state into two groups. First, the methods that require to be configured using expert knowledge. These methods transform raw observations into state representation, however, the hyper-parameters optimization and finding proper settings are left for the experts. Second, the methods that automatically transform the observations to states and reduce the role of expert.

2.1.1. Transferring raw observations to state representation

Manipulating raw observation and constructing new features are widely used for deriving state representation. These methods range from simple approaches such as tile coding applied to linear function approximation methods for problems like n-state random walk [67, 1], to more complex methods like structure2vec [17] and Pointer Networks [78] for graph Combinatorial Optimization Problems (COPs) such as Vertex Cover [44] and TSP [8]. They are mainly employed for expanding the observation to more useful representations with or without taking the final policy into account. For problems like combinatorial optimization [50], robot navigation [2] and real world business problems like train shunting [56] and online advertising [62], raw observation of the environment may require processes to derive state representation.

Each state consists of a n-dimensional observations vector from the environment and each observation is a scalar value called a feature. In other words, the observations in an n-dimensional observation space are feature vectors with n

entries. In more complex tasks like image classification where the observation is a matrix rather than a vector, they can be flattened to vectors. Exact observations from the environment in some problems are not sufficient for representing states. As an example, assume a task with 2-dimensional observation space where the two dimensions are interrelated. In other words, possible actions for the situation that both dimensions are positive or negative are different from when they have different signs. In this case, using a vector with two entries as a state does not take any interaction between the two variables into account. Since, the observations in many environments are represented by numerical values, features can be interpolated to generate new meaningful features. One of the simplest families of features used for interpolation is *Polynomial Features* [74]. Polynomial features are obtained by modeling the observations as any order-n polynomial. Formally speaking, assume $O = (o_1, o_2, ..., o_k)$ is an observation vector from the environment. The new state s_i corresponding to an observation o_i is defined as:

$$s_i = \prod_{j=1}^{k} o_j^{c_{i,j}} \tag{1}$$

where, $c_{i,j}$ is an integer denoting the degree of j^{th} term of observation in the definition of the i^{th} term of state representation. This approach is mainly used for deriving state representation of linear function approximation algorithms when the important interactions between the features are not included in the observations of the environment.

Coarse Coding is another useful approach for generating features especially when the observation of the environment is not informative enough [73]. For example, assume a task with two dimensional state space where each region of the space has its own characteristics. In order to capture pertinent information about the environment, coarse coding introduces some overlapping circles whose status shows the corresponding state of observation. Each observation lies in one or more circles and circles are called present/absent or active/inactive based on the location of the observation. If the observation lies in i^{th} circle, its

Figure 2: An example of Coarse coding. The resulting feature vector of state s is $(0, 0, 1, 0, 0, 1)$.

corresponding value in state s is 1, otherwise it is 0. Using this method, the feature vector of states extends to n binary values. Figure 2 shows an example of coarse coding with $n = 6$.

Tile Coding [74] is a widely used approach for converting continuous space to discrete which is easier to manage and reduces the complexity of the problem. In tile coding, n *tiling* that each has a fixed number of *tiles* are offset from each other by a uniform amount in each direction. Figure 3 shows an example of tile coding.

With the popularity of Deep Neural Networks (DNNs) in past few years, the feature engineering is mainly performed by DNNs and researchers focus more on designing the DNNs. Nevertheless, for some problems like COPs, processing the raw representation of problem instances in order to derive effective state representation improves the quality of the solution. For this reason, many research focus on developing methods for transforming the original observation of the environment to state representation. In [17], an approach named structure2vec for representing structured data like trees and graphs is introduced. This approach is based on the idea of embedding latent variable models into feature space. A vector for representing graph is obtained by employing probabilistic kernels to find latent variable models and a neural network is trained to output the embedding of a graph based on nodes' attributes. This idea is

9

Figure 3: An example of Tile coding. The active tiles are shown in bold margins. The oval is the original observation space and the point is a sample observation.

used in [44] for solving graph-based combinatorial optimization problems such as Minimum Vertex Cover and TSP. The graph embedding network is learned by fitted Q-learning and the output of the network is used as a greedy policy to incrementally create the solution of the problem.

2.1.2. Automatically defining the state representation

There are some approaches in the literature that try to automatically find state representation. Adaptive Tile Coding [83] is one of them where the algorithm starts with a single big tile and two heuristic algorithms determine when and where to split the tiles, respectively. First heuristic keeps tracks of lowest Bellman error and whenever it fails to change, it is the time for splitting. The second heuristic then checks the states in a tile and splits the tile if the number of conflicts in action selection is higher than a threshold. Since the tiles are split evenly, in [48], genetic algorithm is used for solving this issue and automatically deriving tile coding and state abstraction of large state spaces. This approach also starts from a single tile containing entire state space and genetic algorithm determines when and where to split the tiles while the splits can divide unevenly. In this modeling, the population contains different tilings that are evaluated by an RL algorithm. Each tiling is demonstrated by a binary decision tree and its

fitness is the performance of RL algorithm using the tiling for state representation. Mutation either shifts the tiles or divide them. Shift operators move the position of splits and divide operators introduce new splits. This approach is applied on Mountain Car and Pole Balancing as two well-known continuous control RL problems [48].

One approach of handling large state space is to reduce the complexity of the state space by aggregating those states that have similar characteristics. In [7], state aggregation for Q learning algorithm in continuous state space is studied. The proposed method combines Growing Neural Gas (GNG) and Q-learning to automatically learn state aggregation during training. The GNG is an unsupervised method that represents the topology of an environment by a small set of units. It adds units and updates the approximation until deriving a representation. In order to aggregate states, a set of codeword vectors are defined to represent states and a region is assigned to each of them based on nearest neighbor quantizer. Then GNG finds the final codeword vectors for representing states by moving the initial vectors or adding new ones. As another example, an extension of tile coding is introduced in [63] to automatically transform continuous state space to discrete and use the resulting state representation for solving knapsack problem. In this method, a single n-dimensional tiling is used to convert numeric item values to discrete. Each item of knapsack problem is considered as a dimension and RL is used to automatically learn the number of tiles in each dimension. In this modeling, states are items and the action is the number of tiles.

Focusing on graph-based state representation leads to other promising works. In [57], a representation named *DeepWalk* is introduced that generates a sequence of nodes using random walk. This model is used in natural language processing and trains a model similar to word2vec for representing a graph. In [86], an asymmetric proximity preserving graph embedding based on random walk is developed that can be used to represent a graph for social network problems like link prediction.

11

2.1.3. Challenges

Approaches for transforming observations to state representation have parameters and settings that properly tuning them would significantly increase the total reward. For instance, although tile coding and coarse coding are useful approaches for handling large and continuous state space, the number of tilings, tiles and circles have to be determined by the system designer or expert. Furthermore, promising NN-based methods like Pointer Network require appropriate information about the problem instance which is still the task of RL designer. Therefore, completely replacing expert knowledge with automation levels is a challenging task in defining state representation.

The other main challenge is generalization of proposed methods. An automated method such as structure2vec is a powerful way to represent graphs. This embedding is only useful for graph-based problems, while a large fraction of RL problems are not inherently graph-based. The same issue holds for Adaptive Tile coding which require special adaptations for a particular task. Deriving a generic state representation methods would improve an AutoRL pipeline drastically which is not well studied in the literature.

2.2. Methods for Automating Actions

In many RL tasks, actions are mainly the decisions of the agent that alter the state of the environment. Different types of actions such as continuous, discrete, multi-dimensional, bounded or unbounded would entail policies with different qualities. For example, a continuous action like prices in a dynamic pricing task can be modeled as either continuous or discrete action space. On one hand, continuous actions might be more precise, however, it is not possible to model them with tabular reinforcement learning or function approximation approaches like DQN that consider an output for each action. On the other hand, although discrete actions are easier to model, modeling a continuous space as discrete might be tricky especially when small changes in the action would have large impact on the total reward. Therefore, deriving a proper action representation is very important as it is difficult to find the best representation for actions that

ends up with the best policy. For this reason, automating the definition of action spaces are necessary for many tasks. Action spaces could be a combination of discrete and continuous for multi-dimensional spaces like robot joints in robot navigation problems. For many continuous control task such as pendulum or BipedalWalker [11], the continuous action space could be discretized to represent discrete actions. In this subsection, we review learning actions and discretizing continuous action space separately.

2.2.1. Learning actions

Action representation learning in order to improve the action values and policy has become popular in recent years. In [76], the action representation of multi-dimensional action spaces is learned using hyper-graph. Hyper-graph is the generalization of graph in which each single hyper-edge could contain one or more vertices. In this modeling, the actions are modeled as vertices in hyper-graph and the goal is to learn the representation of hyper-edges. To achieve this, a parametric function is defined for each hyper-edge whose input is a representation of the states and the outputs are separate values for each possible action in the action space. If the action space is multi-dimensional, the number of outputs is equal to the cardinality of a Cartesian product of action vertices in the hyper-edge. After receiving a state representation, each parametric function corresponding to a hyper-edge returns a vector for each action and these vectors are mixed up using a non-parametric fixed function. The output of this mixing function is the Q value for RL. In order to find a good hyper-graph for this problem, a rank r is chosen that defines the set of all hyper-edges with the cardinality of at most r. The desired hyper-graph is the one that r is equal to the number of vertices.

An efficient way to represent actions is to model the output of the policy network as a continuous Probability Density Function (PDF). In common practice, Gaussian distribution is used for the policy and mean and standard deviation are learned during the training. Gaussian distribution has been successful in many tasks with continuous action space, however, infinite support of this distribu-

tion might introduce bias in policies obtained from policy gradient algorithms. To solve this issue, Gaussian distribution is replaced by Beta distribution in [15]. The authors show that using Beta PDF for policy reduces bias whilst the performance is not negatively affected.

One approach to automatically derive the policy distribution is introduced in [77]. According to this work, the policy gradient updating rule with parametric distribution functions results in sub-optimal policies. This sub-optimality is in distribution space and learning the policy distribution is a solution. For this purpose, Distributional Policy Optimization (DPO) as an updating rule is presented that minimizes the distance between the policy and a target distribution. This updating rule is shown in Equation 2.

$$\pi_{k+1} = \Gamma(\pi_k - \alpha_k \nabla_\pi d(\mathcal{D}_{\mathcal{I}^{\pi_\parallel}}^{\pi_\parallel}, \pi)|_\pi = \pi_k) \tag{2}$$

where Γ is a projection operator onto the set of distributions, d is a distance measure, $\mathcal{D}_{\mathcal{I}^\pi}^\pi$ is a distribution over all states and actions that their advantage value is positive. In order to minimize the distance between two distributions, Implicit Quantile Network [16] is employed by using Wasserstein distance metric. DPO algorithm is incorporated in Generative Actor Critic that contains four neural networks for policy, delayed policy, value network and action-value network.

2.2.2. Discretizing continuous actions

Continuous action spaces are challenging in many tasks. As mentioned before, some RL algorithms like DQN do not work well for continuous action spaces because they rely on ϵ-greedy algorithm and the best action is required at each step. Finding the best action in a continuous space needs an optimization step for each interaction with the environment which is intractable. For discrete action spaces, a separate output is considered in the policy network for each action which is not possible when the action space is continuous as it needs infinite number of outputs. Although discretization [42] makes the action space

discrete and manageable, it is not suitable for tasks that are very sensitive to small alteration of actions.

Sometimes, continuous action spaces could be transformed to discrete while retaining necessary information for action selection as shown in [75] for on-policy control RL when the domain of all the continuous actions is between -1 and 1. The set of discrete actions for each dimension is $\mathcal{A}_i = \{\frac{2j}{K-1} - 1\}_{j=0}^{K-1}$, where, K is the number of discrete actions. The discrete policy is a neural network that outputs a logit L_{ij} for j^{th} action in i^{th} dimension. For each dimension i, the logits are combined through a softmax function to compute the probability of choosing each action. This approach is integrated with TRPO and PPO to be evaluated on MuJoCo benchmarks [75].

2.2.3. Challenges

Similar to state representation, approaches for determining action space contains hyper-parameters and finding appropriate configurations are important. One way to optimize the parameters is through hyper-parameters optimization module as it is depicted by a dashed arrow between actions and hyper-parameter optimization in Figure 1. This is challenging because there might be several action spaces and each has some parameters. Hence, the search space is relatively big and the optimization procedure is computationally expensive.

Finding the optimal action space needs many trial and error steps by checking different definitions. Automation level may help to decide the policy distribution or descretization approach for continuous actions is necessary. Designing the structure of the policy output is normally performed using expert knowledge which is not always available. This is an interesting research direction that may influence positively on the output of a RL framework.

2.3. Automated Reward Function

Reward function is the other component of MDP that greatly influences on the quality of the policy. Assume that an agent aims to reach a goal by searching in a 2-D grid world. Different reward functions lead to different policies. If all

moves have zero reward except the one that takes the agent to the goal state, the obtained policy is totally different with when the reward pertains to the distance to the goal. In the former case, the agent might waste time by wandering in the grid field since there is no penalty for the moves. Therefore, designing a strong reward function is of great importance which is normally the task of RL experts. Even using expert knowledge, sometimes it is necessary to test different reward functions in order to obtain the one yielding the best policy. Hence, automating reward function can largely help the agent toward finding a good policy. Generally, there are three approaches for designing reward function that can be considered in an automation level. An agent may automatically seek among these approaches to find the best representation for the reward function.

2.3.1. Curriculum Learning

Curriculum learning is a useful method for learning sparse reward functions. Many problems like robots navigation have sparse reward function because the search space is extremely large and only goal state produces a positive reward. Curriculum learning gradually increases the difficulty of the problem by starting from the states close to the goal state and progressively increase the distance. A curriculum learning method inspired by dynamic programming methods is introduced in [27] for goal oriented tasks like robot navigation. The motivation is to handle the sparsity of the reward function because the reward value is zero for all the moves except the one that takes the agent to the goal state. Based on this work, the robot starts to learn by searching from states that are close to the goal. The process is continued until showing mastery on the current setting. Then, new start states are added to the search space by a random walk from previous start states. This so called *reverse learning* is helpful to solve the problem of sparsity.

Focusing on the limitation of the reverse learning with random walk, a new approach is presented in [38] for highly dynamic and unstable systems where previous method breaks down. Similar to [27], the backward learning method in [38] starts from the state in the vicinity of the goal state and increases the

distance when the agent demonstrates mastery in solving the problem. Unlike [27], where the state space expands by random walk, in [38] the new states are obtained by computing approximate backward reachable set (BRS). BRS is the set of all points in the state space that the agent is able to reach a certain region in a fixed and short amount of time. The main contribution of this work is to update the state space in the way that only reachable states from goal state are preserved.

Although curriculum learning is helpful in solving sparsity, their main application is in goal searching tasks. Using curriculum learning in COPs is quite challenging. Learning heuristics in these problems might not be adapted with curriculum learning because the optimal solutions are not usually available in advance. Therefore, using curriculum learning in a reinforcement learning pipeline is limited to goal-based problems where the goal state is available.

2.3.2. Bootstrapping

Bootstrapping methods start learning from a pre-defined policy. This policy could be either for a similar task or designed by human. A typical approach based on bootstrapping is introduced in [69] where the learning process is split into two phases. In the first phase, the robot is controlled by a supplied control policy or directly by human. In the latter case the robot is navigated by human in the environment and it updates the value function during this phase without changing the policy. The second phase is a typical reinforcement learning process and the robot updates the policy based on the value function which is initialized in the first phase.

In order to address three main RL challenges including sparsity of the reward function, lack of effective exploration and brittle convergence properties, a hyper algorithm that incorporates evolutionary algorithms (EA) and DRL is introduced in [43]. Each individual in the evolutionary algorithms represents a DNN and mutation randomly alters the weights of a DNN. These DDNs are the actors in Deep Deterministic Policy Gradient (DDPG) [47] algorithm. Initially, a population of actors are generated and the fitness for each actor is the

17

cumulative reward of following its policy for a particular episode. Then the actors are perturbed through mutation and crossover to create new set of actors as the next generation. Through crossover, an offspring policy is obtained by randomly exchanging segments of parents' weights. The mutation is performed by applying random Gaussian noise to the weights. The experiences of actors are stored in a replay memory and these experiences are used by another actor (apart from individuals in EA) to update its weight through DDPG algorithm and this actor is copied to the EA population periodically. The algorithm is similar to DDPG except that the main actor has access to the experiences of individuals. This approach is categorized as bootstrapping as it uses information of other policies for training the target policy.

Since DNNs are usually sensitive to small modification of the weights, mutation may cause extent forgetting. To address this issue of [43], the evolutionary RL algorithm is modified in [10] by adding personal replay memories and defining a crossover and a mutation operators. A small personal replay memory named genetic memory is defined for each individual to keep their most recent experiences. In the new crossover named Q-filtered distillation, the genetic memory of the child policy is filled with the latest experiences of its parents. The child policy is initialized with the weights of one of its parents and it is trained using the experiences of its genetic memory. The action selection procedure is toward minimizing a loss function that combines the policies of the parents. The new mutation named proximal mutation modifies the Gaussian perturbation of [43] by scaling the Gaussian using the sum of gradients over a batch of transitions. This approach is evaluated on MuJoCo environments including Hopper, Ant and HalfCheetah.

One challenge in sparse reward tasks with high-dimensional state space is deriving a policy for long horizon. On one hand, while planning algorithms work well for task that determining a distance between states and designing local policy is easy, it is problematic when it comes to high-dimensional observations. On the other hand, even though RL algorithms work well with high-dimensional observations, they fail to solve some tasks in long horizons. To achieve the

benefit of planning and RL algorithms, a combined algorithm is presented in [23]. In this approach, the task is decomposed into a series of easier goal-reaching tasks and the problem of finding these sub-goals is reduced to a shortest path problem using a distance metric. By defining the reward value as -1 for every move until reaching a goal, the Q function and the shortest path would be closely connected. Using the distance between nodes, a directed graph is built on top of the states in a replay buffer in which each node corresponds to an observation. This graphs finds a set of nodes showing the shortest path from the start state to the goal state and the goal-based policy starts to find the goal by reaching intermediate goals one after each other.

2.3.3. Reward Shaping

Reward shaping is the process of learning a proxy reward function that maximizes the expected return. Although rewards in many tasks like video games are straightforward to define, it needs careful considerations in problems like robot navigation and COPs that have multiple objectives. For example, finding a reward function by combining several objectives like avoiding obstacles, reaching the goal, passing minimum distance and the optimal angular velocity in a goal searching task is challenging. As another example, the optimal solution of a TSP with time window and prizes should has the minimum tour length and maximum collected prize while the time window constraints are satisfied [37]. For these kinds of problems, different rewards may lead to completely different solution performance and carefully designed rewards are necessary. In the next paragraphs of this sub section some example and application of reward shaping in different domains are reviewed.

In [41], reward shaping in the context of imitation learning is explored. In imitation learning, the agent faces an MDP without any reward value. States and actions pairs from a target policy are demonstrated to the agent and the goal is to mimic the target policy using these state-action pairs. In [41], a reward value is provided for the agent which is not necessarily aligned with the target policy. The accuracy of the learned policy is a parametric function where the

19

parameters are obtained by maximizing the shaped reward.

In [72], the benefit of reward shaping in Spoken Dialogue Systems (SDS) as a Partially Observable MDP (POMDP) is studied. In SDS, an agent interacts and converses with the clients and a numerical reward value is received at the end of the conversation. Since there is no intermediate reward, the reward function is sparse. In this work, domain knowledge is used to provide another reward signal. This extra reward value is obtained from an RNN which is trained by supervised learning using previous annotated data. The state (belief in POMDP context) and action pairs are the inputs of the RNN and the scalar return of the dialogue is the target value. The agent uses the extra reward value along with the feedback of the environment to learn its policy.

A well-known reward shaping approach is potential-based reward shaping [54]. Let $r(s, a, s')$ be the instant reward of taking action a in state s and going to state s'. Potential-based reward shaping employs a function $F(s, a, s') = \gamma \Phi(s') - \Phi(s)$ where, γ is discount factor, and $\Phi(s)$ and $\Phi(s')$ are potential functions of states s and s', respectively. The new reward function is $r(s, a, s') + F(s, a, s')$ which is the sum of original reward and the potential-based reward shaping. In [31], potential-based reward shaping is analyzed in episodic RL. The authors show that temporal credit assignment or reward value could speed up training deep neural networks if a potential based reward function is added to the original reward. An important requirement of this work is that the policy derived by shaped reward should be equivalent to the original policy. Potential-based Reward shaping preserves the optimal policy when the goal states are predefined terminal states and shaping reward is zero in the goal states [54]. It is shown in [31] that policy invariance is violated for finite horizon domains with multiple terminal states and they propose to set the potential value of terminal states as zero to solve this issue.

Learning a potential function $\Phi(s)$ using Meta-Learning in order to accelerate the learning is studied in [89]. The potential function is defined as a neural network and its parameters are updated using Model-Agnostic Meta-Learning

(MAML) algorithm [26]. The updating equations of MAML are as follows.

$$\phi_i = \theta - \alpha \nabla_\theta \mathcal{L}_{\tau_i}(f_\theta) \tag{3}$$

$$\theta = \theta - \beta \nabla_\theta \mathbb{E}_{\tau_i} \mathcal{L}_{\tau_i}(f_{\phi_i}) \tag{4}$$

where, τ_i is a MDP task deriving from a particular task distribution, θ is the parameters of the value network, ϕ_i is the parameters of potential-based reward shaping which is task-specific, α and β are learning rates, $f(\theta)$ and $f(\phi_i)$ are the value network and reward shaping network respectively, and $\mathcal{L}$ is the loss function. The parameters of the networks are trained through running an adapted version of DQN with replay memory.

The term *AutoRL* is used in [14] to address automatically generating a reward function for a robot navigation task. This work defines the learned reward functions for path following and target-based tasks in robot navigation. The motivation of the proxy reward function is to handle the sparsity of the original reward which is a binary value showing whether the agent reaches its objective. More specifically, the objective of the target-based task is maximizing the probability of reaching a goal as shown in Equation (5).

$$G_{P2P}(s) = \mathbb{I}(||s - s_g|| < d_{P2P}) \tag{5}$$

where, $P2P$ shows the *peer to peer* property of the problem, G is the objective function and $\mathbb{I}$ is the indicator function to determine whether the distance between current state s and goal state s_g is less than d_{P2P}. Using this objective function, the reward is defined in Equation (6).

$$R_{\theta_{r_{P2P}}} = \theta_{r_{P2P}}^T [r_{step} \ r_{goalDist} \ r_{collision} \ r_{turning} \ r_{clearance} \ r_{goal}] \tag{6}$$

In this formulation, r_{step} is a constant penalty for each step, $r_{goalDist}$ is the negative Euclidean distance to the goal, $r_{collision}$ is a binary value showing whether the robot collides an obstacle, $r_{turning}$ is the negative angular velocity, $r_{clearance}$ is the distance to the closest obstacle, and r_{goal} is a binary value

showing whether the goal is reached. Using this reward function, an actor-critic algorithm is employed to learn the optimal policy. The algorithm consists of two steps: learning the parameters of the reward function and training actor-critic networks. In order to learn the parameters of the reward, they are initialized using previous evaluated parameters according to a black-box gradient-free optimization algorithm. Then the parameters of the fixed-shape actor-critic networks are learned using the reward function. Upon running the algorithm for a certain amount of time, the parameters providing the maximum objective value are selected for reward function and RL trains actor-critic networks.

In [24] the idea of learning a parametric function is employed for a set of typical continuous control RL problem like *Ant*, *Walker2D*, *HumanoidStandup* and *Humanoid*. These environments are implemented in *OpenAI Gym* [12]. For each problem, a particular parametric reward function is defined and the same algorithm is used for learning both the parameters of reward and policy network. Actor-critic algorithms including Proximal Policy Optimization [65] and Soft Actor Critic [34] are used with parametric reward function and the method outperforms the same algorithms without parametric reward on the aforementioned tasks.

A reward vector based on limited feedback of environment is presented in [61] for optimizing reserve prices in real time auctions. In this work, the reserve price of ad slots in online adverting is the action. An agent sets the reserve price using a policy network whose input is the ad slot information and its output is the a probability distribution over reserve prices. The environment returns only a binary value determining whether the ad slot is sold in the auction. This binary feedback is not sufficient for learning a policy for reserve price setting. In order to improve the reward function, the valid interval of reserve prices is divided into some sub-interval and a weight is assigned to each sub-interval. The reward value in this case is the inner product of the weight vector and a reward vector. The reward vector has a separate entry for each interval and activates each entry only if the reserve price lies in corresponding interval. Using this method on a simulation showed improvement in the revenue.

2.3.4. Challenges

As mentioned before, Curriculum learning is appropriate for goal-searching problems. This method fails for path following problems where the goal state is unknown. Using this method in AutoRL is challenging because the goal state is mainly the solution of the optimization problems. However, the solution of easy instances of a problem might help to solve complex instances, and this generalization is relatively hard but important progress in AutoRL.

Bootstrapping methods require initial policies that are provided either by expert knowledge or by following other learning or optimization methods. Though expert knowledge are helpful, they are not available for most of problems. Besides, using any policy other than optimal policy do not help the learning because that biases the value function and a sub-optimal policy is learnedNevertheless, finding the optimal approach is time consuming and a level of automation could be largely beneficial.

Reward shaping methods have some parameters that are tuned before starting the training phase. Similar to parametric state and action representation methods, these hyper-parameters settings highly influence the total reward. Tuning the hyper-parameters is challenging and time-consuming which makes it an interesting research direction. Another interesting research direction is about the effects of intrinsic rewards whether they help in solving the problem of sparse rewards [68]. For example giving a reward when the agent has a lot of options available, or when the agent encounters a new situation.

3. Automated Algorithm Selection

When the problem is modeled as a sequential decision making problem and the components of MDP are defined, the next step in solving it with RL is to select an appropriate algorithm. One way to reduce the search space is to filter the algorithms based on the class of the problem. For example, if the states and actions are discrete and finite, tabular RL algorithms like typical Q-Learning and SARSA [74] are suitable candidates and there is no need to search over the

23

class of sophisticated algorithms in Deep Reinforcement Learning. Moreover, if the model of the environment is known, a wide variety of model-based algorithm like dynamic programming could be utilized. Despite having different context, algorithm selection approaches in AutoML can provide insights for AutoRL [60]. Most of work in the domain of RL algorithm selection are intertwined with hyper-parameter optimization. For this reason, we explain these combined algorithm selection and hyper-parameter optimization in the next section and present few work that merely focus on algorithm selection in this section.

In [19], algorithm selection in supervised learning is modeled as a contextual multi-armed bandit problem. This approach is developed for AutoML. The difference with RL is in using a context vector that contains the information of the dataset. Each decision moment starts with observing the dataset and its feature vector which is the context vector in contextual multi-armed bandit. Then Upper Confidence Bound (UCB) and ϵ-greedy algorithms are used for learning the values of arms on a set of datasets.

The algorithm selection problem in episodic RL tasks is modeled as a multi-armed bandit problem in [45] to decide which RL algorithm is in control for each episode. A set of RL algorithms are given and the process of algorithms selection is started with an empty trajectory set. At each time step, an algorithm is selected and it generates a trajectory with discounted reward according to the policy of the selected algorithm. RL algorithms are selected based on Epochal Stochastic Bandit Algorithm Selection in which the time-scale is divided into epochs of exponential length. The policies of algorithms are only updated at the start of epochs and they are not changed during the epochs. This way of updating handles non-stationarity induced by the algorithm learning.

3.1. Challenges

Generally speaking, the RL algorithm can be classified according to the type of the policy or value function. Through this categorization, the policy or value function could be either tabular or parametric. One initial challenge in RL algorithm selection is to decide between these two. If the state and action spaces are

relatively small, tabular methods are more appropriate, whereas these methods do not work for large and continuous space and action spaces. Discretization is another solution that is discussed in section 2.2.

After identifying appropriate class of RL algorithm, selecting an algorithm to learn the policy is another challenge. One way is to treat the algorithm as a hyper-parameter and optimize that in the hyper-parameters optimization module which is normally followed in AutoML. However, this approach require particular hyper-parameters optimization framework because the quality of the algorithm depends on the problem, its MDP modeling and parameters settings. Furthermore, the number of required timesteps varies for different algorithm and this makes comparing the algorithms challenging. In sum, selecting proper RL algorithm for a task is difficult and it highly depends on the problem.

4. Hyper-Parameter Optimization

An optimal RL configuration for solving a sequential decision making problem highly depends on promising hyper-parameters settings. Hyper-parameters are fixed during the training and they are usually set by RL experts prior to starting the interactions. For example, learning rate in policy gradient or value function updating formula, discount factor, eligibility trace coefficient and parameters of a parametric reward shaping method are hyper-parameters. Different tasks require different sets of hyper-parameters which make hyper-parameters optimization challenging and automating this process would be very useful. Many different approaches are developed for automatically optimizing hyper-parameters of supervised learning algorithms that they could be adapted with RL to optimize hyper-parameters of RL algorithms. In this section, we first review hyper-parameters optimization approaches and then their applications. At last, the main challenges of optimizing hyper-parameters are elaborated.

4.1. Methods for tuning hyper-parameters

This subsection presents the previous hyper-parameters optimization work categorized by their core methodology.

Stochastic Gradient Descent. Backpropagation is a main method for training neural networks in which the gradient of the loss function is computed with respect to the weights. This gradient is propagated backward through the network and the new weights are obtained by a variant of gradient descent algorithm. The hyper-parameters of gradient descent with momentum including decay rate and learning rate are included in backpropagation algorithm and they are optimized together with neural network weights in [49].

Bayesian Optimization. Bayesian Optimization methods including Sequential Model-based Algorithm Configuration (SMAC) has been very popular in AutoML [36]. These methods are beneficial for optimizing expensive to evaluate functions such as the performance of supervised learning algorithms. The idea of Bayesian Optimization is extended to RL hyper-parameters optimization in [6]. In this work, *RLOpt framework* uses Gaussian process regression as surrogate function and integrates Bayesian optimization with RL. The process of hyper-parameters optimization is considered as a supervised learning problem where the parameters are the input and the performance is the target. At each step, the selected parameters are given to the agent and it learns a policy and returns the performance. Based on the history of (hyper-parameter, performance) tuples, Gaussian Process is used to select the next hyper-parameters settings.

Multi-Armed Bandit. A bandit-based hyper-parameter optimization algorithm named *Hyperband* is proposed in [46]. This algorithm is based on *Successive-Halving* [39], where a space of n hyper-parameters is searched. At each iteration, all the hyper-parameters configurations are evaluated and the worst half are removed from further processing. This continues until only one configuration remains. Unlike SuccessiveHalving, in Hyperband, the value of n is not predetermined. This value is set according to the total processing budget and the maximum amount of resources that can be allocated to a particular allocation. In each iteration, n is determined and then SuccessiveHalving is invoked to find the best hyper-parameters.

Evolutionary Algorithms. In [25], the parameters of $SARSA(\lambda)$ and $Q(\lambda)$ as two RL algorithms based on eligibility traces are optimized using Genetic Algorithm (GA). In this method, a vector containing all hyper-parameters is considered as a *chromosome* and the mutation and cross-over are performed on this vector. The algorithm is tested on under-actuated pendulum swing-up and the authors show that the selected parameters maximize the end performance. Another application of GA for optimizing the hyper-parameters of RL algorithms is presented in [66], where the parameters of Deep Deterministic Policy Gradient (DDPG) with Hindsight Experience Replay (HER) [4] are learned through GA. The target parameters for GA are discount factor, polyak-averaging coefficient which is used for updating target networks in algorithms like DQN with different target and main networks, learning rate of actor and critic networks, percentage of time a random action is taken and Gaussian noise parameters. Concatenation of binary representation of these parameters builds the chromosomes and the fitness is the inverse of number of epochs to reach close to maximum success rate for a particular task.

Normally, there are three challenges in hyper-parameters optimization of DRL. First, dynamic environments require dynamic hyper-parameters and optimal hyper-parameters settings in one stage might poorly work in another stage. Second, the optimization is not sample efficient and needs a full training run to test each selected hyper-parameter. Third, dynamic modification of the neural network is not considered in literature. A joint optimization approach based on evolutionary algorithms that optimizes the agent's network and its hyper-parameters simultaneously is presented in [29]. In the evolutionary framework, each individual is a DRL agent consisting of a policy and a value network together with RL algorithm paramteres. Rollouts of each individual is stored in a shared replay memory to be used as experiences for other agents. Each agent is evaluated by running for at least one episode in the environment and the mean reward of the agent is used as its fitness. After crossover and mutation, all the agents in the environment are trained using the experiences in the shared replay memory. This approach is applied on TD3 algorithm in MoJuCo contin-

uous control benchmark. This integration of evolutionary algorithms and neural networks is known as neuroevolution [71].

Greedy Algorithms. In order to optimize the decay rate in algorithms based on eligibility trace like $TD(\lambda)$, a greedy algorithm is proposed in [82]. This algorithm defines λ as a function of states for RL algorithms with linear function approximation. In each iteration of policy evaluation algorithm, the agent takes the value of λ greedily according to the weight vector, observation vector corresponding to the current and the next states, instant reward and importance sampling. The intuition of this greedy algorithm is to minimize the error function which is the difference between the return obtained from the selected λ and the Monte Carlo return ($\lambda = 1$).

Reinforcement Learning. Hyper-parameters optimization is modeled as a sequential decision making problem in [40] and RL is used to find the optimal hyper-parameters. In this framework, the agent learns to explore the space of hyper-parameters of a supervised learning algorithm and the final parameters minimize error on the validation set. This method is naturally for supervised learning and works based on training and validation datasets. However, the general idea can be extended to RL algorithms. In the MDP modeling of the hyper-parameters optimization problem, the state of the environment is defined as the meta-features of an input dataset plus the history of evaluated hyper-parameters together with their performance, the action is a value for each hyper-parameter and the reward is the performance of ML algorithm on the input dataset with the selected hyper-parameters. Upon modeling the problem as MDP, deep Q network algorithm [53] is used for learning the parameters of an LSTM system and the obtained policy works as a decision maker to determine optimal hyper-parameters for each dataset.

In common practice, the hyper-parameters of an algorithm are optimized once and they are fixed during entire run of the algorithm. However, because most of AI algorithms are iterative, the optimal hyper-parameters might change over time. The problem of dynamic algorithm configuration is studied in [9]

and RL is used to derive a policy for optimal configuration in each step. In this modeling, states are descriptions of an algorithm A and actions are assigning particular values to hyper-parameters of A. Reward function depends on the instances drawn from the same contextual MDP. The optimal policy is obtained either by a tabular Q-learning or DQN to select the configuration that has the highest discounted reward. Since instances' information are part of state, the optimal configuration might be different for different instances.

Neural Networks. One main challenge of using well known hyper-parameter optimization in AutoML such as Sequential Model-Based Optimization (SMBO) and SMAC, is the time needed for performing necessary iterations. These iterative algorithms take remarkable amount of time to optimize the hyper-parameters and this is highly prohibitive in RL framework. This challenge is the motivation of developing a neural network for finding a mapping between data and hyper-parameters [13]. The meta-features of the dataset are the input of a Convolutional Neural Network (CNN) and hyper-parameters of the algorithm are the output. Training the CNN is based on supervised learning using subsets of a large dataset as the training data. The target hyper-parameters used for supervised learning are obtained by Bayesian Optimization.

4.2. Applications

In [20], hyper-parameters optimization for object tracking algorithms is modeled as RL and Normalized Advantage Functions [32] is used to learn a policy network that receives a state and returns the optimal hyper-parameters settings for a particular object tracking algorithm. Most of these algorithms produce a heat map in the search region showing the probable location of objects. The combination of this heat map, the parameters of the object tracking algorithm and the appearance features like RGB-color constitute the state, and the reward is the tracking accuracy.

One important research direction in robotics is domain randomization. Domain randomization is a technique to use a simulation model to provide a policy

for real environment. In other words, when training in real environment is not practical, domain randomization helps to derive a policy that maximizes the total expected return over a set of MDP obtained from the same distribution. The parameters of this MDP distribution are usually fixed and predefined, however, fixed parameters might not be sufficient for some environments. For this reason, Cross Entropy Method is used to learn these parameters in [79]. In this method, the policy parameters function is a function of MDP distribution parameters and optimal policy parameters are derived by PPO. The MDP parameters are acquired by maximizing discounted return of following the optimal policy in real environment.

Application of RL in different domains requires special consideration for tuning the hyper-parameters. For instance, in [55], RL is leveraged to solve Sequential Ordering Problem (SOP) - a variant of TSP with a precedence constraint. Tuning parameters in this work are performed by testing different RL algorithms including SARSA and Q-learning, different reward definition and some different values for ϵ in ϵ-greedy. This configuration of the algorithm is used to solve SOP.

4.3. Challenges

Although several methods are developed to tune hyper-parameters, hyper-parameters optimization for RL algorithms may be computationally expensive because the agent needs to interact with an environment and update its policy continuously for a number of timesteps or until reaching convergence. This is usually a time-consuming and intractable process that needs special considerations. Efficient methods in AutoML such as Bayesian Optimization may work well for RL models with small state, action and trajectories. However, for complex tasks with large state and action space or long trajectories, existing methods require considerable adaptation to provide the best configuration in a reasonable time. This is a big challenge in RL hyper-parameters optimization.

One main difference between supervised learning and reinforcement learning is the evaluation criteria. Since, there is no target value like supervised learn-

ing datasets to show the desired behavior. Instead, an agent seeks to find a compromise between exploration and exploitation that helps as hints [6].

Multi-fidelity methods [28] does not run the ML for the full budget for every hyper-parameters but only for a limited budget (low fidelity). Then, only promising hyper-parameters are run for longer (high fidelity). Adapting this method to AutoRL is challenging and also interesting because it might reduce the required time budget for optimization.

5. Learning to Learn

Apart from the three main components of a RL framework, levels of automation are also presented in the literature for the procedures that can be placed in more than one components. For example, gradient descent method is used for updating the parameters of parametric functions like policy or reward. In this section, recent works on automating these kinds of procedures are reviewed. Most of these works are inherently developed for supervised learning. However, the same motivation and requirements hold for RL which shows interesting research directions for future works.

In [3], the normal gradient descent formula is replaced with a new formula in which a function of gradient value is used rather than the original gradient in the updating rule. The new gradient descent update equation is shown in Equation (7).

$$\theta_{t+1} = \theta_t + g_t(\nabla f(\theta_t), \phi) \tag{7}$$

where, θ is the parameter of the objective function $f(\theta)$. In this formulation, $g_t(\nabla f(\theta_t), \phi)$ is a function of the gradient of f with parameters ϕ which is obtained by a recurrent neural network (RNN). As shown in Figure 4, the method consists of two neural networks. Function f known as *optimizee* is represented by a feed-forward neural network with parameters θ. The gradients of f are used in function g which is represented by an LSTM recurrent neural network. These gradients plus the hidden states of RNN are the input and g
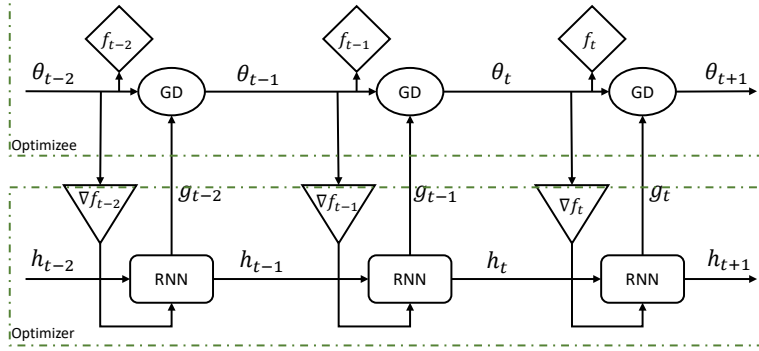
Figure 4: The gradient descent using RNN proposed in [3].

is the output which is used in the updating rule shown in Equation (7). The method is tested on a class of 10-dimensional quadratic function and also on MNIST and CIFAR-10 datasets and the results show the power of using RNN for g.

Selecting RL algorithms is normally performed by expert knowledge or using approaches explained in previous sections. Instead of using existing RL algorithms that perhaps each of them works well on certain type of problems, a model is introduced in [80] that can learn a RL algorithm. Specifically, a distribution D over MDPs is defined and a RL algorithm is learned in the sense that it performs well on the MDPs drawn from D. During solving MDP with reinforcement learning, an RNN is trained that its inputs are states, actions and rewards of the MDP and the output is the policy. Therefore, a recurrent neural network works as a RL algorithm.

Neural Networks are mainly trained using a variant of Stochastic Gradient Descent (SGD) such as normal SGD, SGD with momentum or Adam. The performance of these algorithms depends on the selected learning rate which would be different for different context and applications. An automatic framework based on RL for deriving the best learning rate is proposed in [18]. In this approach, a set of features are introduced to represent the states in the RL modeling. These features include the variance and the gradient of the loss func-

tion. The state representation is used to train a policy using Relative Entropy Policy Search (REPS) algorithm [58] for deciding the learning rate of a particular optimizer. REPS ensures the policy updates to be close to each other by constraining the updates through a bound on Kullback-Leibler (KL) divergence.

A general meta-learning approach that can be applied to any model learned by gradient descent, is presented in [26]. The goal of this approach is to update the model's parameters using few training steps in order to produce acceptable results on a new task. For this purpose, a parametric model is defined that aims to work well on tasks drawn from a task distribution. The algorithm starts with random initialization of model parameters. During each iteration, a set of tasks are sampled from the given distribution and the tasks' adapted parameters are updated using gradient descent on a particular number of examples. At the end of each iteration, the model's parameters are updated using the adapted parameters. The paper discusses the application of this method in supervised learning and classification, and a possible extension to RL is also explained.

In [22] a reset policy is considered together with the reinforcement learning policy to reset the environment prior to an expensive-to-reset situation. For example, an autonomous car would crash in high speed and resetting the environment would be very costly. After the crash state, the environment should be reset manually. Based on this example, a reset policy is necessary for some tasks to decrease the number of manual resets. In [22], off-policy actor critic method is used to learn policies, where the Q values of the main policy and the reset policy are jointly learned. The reset policy takes over selecting an action to abort the episode if its Q value for a particular action taken by forward policy is lower than a threshold. The safe actions are the reversible sequence of actions where the agent can always undo them.

6. Automating Neural Network Architecture

Combining DNNs and RL introduces several successful algorithms in solving complex problems like COPs and video games. Although using DNNs improves

the quality of function approximation, the performance of the DRL algorithms highly depends on proper structure of DNNs. Different methods have been proposed in the literature for automatically defining the best DNN structures. These methods can be categorized as hyper-parameters optimization; however, we assign a separate section to emphasize on their importance.

In [87], a recurrent neural network - the controller - is trained with reinforcement learning where the outputs of this RNN determine the architecture of another neural network - the child network - that is used for prediction. The child network is configured by the controller and it is trained using a dataset. The obtained performance is used as reward for training the controller. This approach is originally developed for supervised learning, although deriving optimal architecture of a neural network can be helpful for other domains.

According to [88], applying the method presented in [87] directly on large datasets is computationally expensive. The solution introduced in [88] is to search on a proxy dataset which is rather small, and then transfer the learned network architecture to a large dataset. The search process is the same as [87] where a RNN provides the architecture of the child network. The search space in this work contains generic convolutional cells that are expressed in terms of repeated motifs in various CNNs like combination of convolutional filter banks. Two types of convolutional cells are introduced with feature maps of the same dimension and half size, respectively. The final CNN architecture is a combination of these cells. The controller receives the output of previous cells and generate the next architecture of the final CNN.

Deriving DNN architecture can be modeled as a sequential decision making problem and RL is a suitable approach for solving that. In [5], a meta-modeling algorithm based on RL named MetaQNN is introduced to generate CNN architecture. The process of CNN architecture selection is automated by a Q-learning agent whose goal is to find the best CNN architecture for a particular machine learning task. The validation accuracy of the given ML dataset is used as the reward value for the agent and the actions are obtained by following ϵ-greedy algorithm and exploring in a discrete and finite space of layer parameters. This

approach shows high performance for image classification tasks.

Optimizing the parameters of NNs is an interesting research area as those parameters greatly influence on the performance. In [85], a DRL approach is presented for automatically learning the value of the learning rate in stochastic gradient descent. Given the model parameters of the neural network and the training samples, the authors use actor-critic policy gradient method to pick a learning rate through the policy network for the gradient descent algorithm. The state in this modeling is a compact vector of the model parameters to avoid processing all the parameters of large networks. Immediate reward for updating the learning rate generator network is the difference between the loss function of the main model in two consecutive time steps. According to the presented results, automatically deriving the learning rate increases the quality of the prediction model.

In [30], the authors focused on the problem of lifelong learning and how an agent learns the optimal policy of a particular MDP using the information of a sequence of MDPs from the same distribution. Specifically, the problem in this work is to search for an optimal exploration policy which an agent follows during exploration in the environment. Each agent maintains two policies: an exploitation policy which is task-specific and an exploration policy which is shared between all the MDPs drawn from the same distribution. At each timestep, each of the two policies provide an action and the selected action is determined by ϵ-greedy algorithm. The exploration policy receives the same reward as the exploitation policy and a variant of policy gradient algorithm (REINFORCE or PPO) is used to update the policy. This approach is experimented on some typical RL problem classes such as Pole Balancing.

Deep neural networks need huge amount of computation for training and this computational complexity is prohibitive sometimes. One approach to reduce the intensity of computation is through quantization of neural networks. Basically, quantization reduces the bitwith of the operations and it can be used to reduce bitwith of layers in neural networks. As accuracy preserving bitwidth may vary across different layers, the problem of learning optimal bitwidth is explored in

[21]. In this work, a DRL approach is proposed to determine the bitwidth of each layer. The states comprises of static information about layers and dynamic information of network structure during RL training. The actions are bitwidth of each layer which are flexible and the agent can change the quantization of each layer from any bitwidth to any other bitwidth. The reward is pertaining to the accuracy and a measure for memory and computation cost. In fact, the two objectives of reward function are preserving accuracy and minimizing bitwidth. Using these definitions of state, action and reward, PPO algorithm is used to learn a policy for deriving the bitwidth of each layer of neural networks.

Components, topologies and hyper-parameters of neural network are automatically determined in [51] using a neuroevolutionary algorithm in which the neural networks are trained using evolution rather than gradient descent. This method that is called DeepNEAT, starts with an initial population of DNNs with minimal complexity and new nodes and edges are added to each chromosome through mutation. Each chromosome contains some nodes and each node is a layer of neural network together with its hyper-parameters. To obtain the fitness of each chromosome, it is converted to a DNN and the DNN is trained for a fixed number of epochs. DeepNEAT is extended to CoDeepNEAT that decomposes complicated DNN structures into multiple repeated modules. Specifically, in CoDeepNEAT there are two populations of modules and blueprints. Each blueprints chromosome contains pointers to a particular module and each module chromosome represents a small DNN. During fitness evaluation, the small DNNs corresponding to the pointers of blueprints are combined to build a large DNN. This approach is evaluated on an image captioning task.

7. Conclusion

In this paper, we presented recent work in automated reinforcement learning that can be included in an automated RL or DRL pipeline. We introduce the general AutoRL pipeline that can be used for solving sequential decision making problem. This area is becoming very popular and a robust and high quality

reinforcement learning pipeline could help in solving many complex tasks and reduce the time and other required resources significantly.

RL framework is split into three main components in this paper and we presented some work related to each component. These three components are MDP modeling, algorithm selection and hyper-parameters optimization. Besides these three, learning to learn methods and learning neural network architecture are also explained in separate sections.

By exploring the relevant literature in AutoRL, we conclude that a concrete and complete pipeline for AutoRL similar to the AutoML pipeline has not been developed yet, even though it has many benefit to design and solve sequential decision making problems. Furthermore, AutoRL is a relatively new research area and it is attracting more and more attention. Optimizing hyper-parameters with minimum required resources, and automatically modeling a problem as MDP and generalizing the mapping from available information to an RL environment are the most important research questions in AutoRL.

8. Acknowledgement

References

[1] Abdoos, M., Mozayani, N., and Bazzan, A. L. (2014). Hierarchical control of traffic signals using q-learning with tile coding. *Applied intelligence*, 40(2):201–213.

[2] ALTUNTAŞ, N., Imal, E., Emanet, N., and Öztürk, C. N. (2016). Reinforcement learning-based mobile robot navigation. *Turkish Journal of Electrical Engineering & Computer Sciences*, 24(3):1747–1767.

[3] Andrychowicz, M., Denil, M., Gomez, S., Hoffman, M. W., Pfau, D., Schaul, T., Shillingford, B., and De Freitas, N. (2016). Learning to learn by gradient

descent by gradient descent. In *Advances in neural information processing systems*, pages 3981–3989.

[4] Andrychowicz, M., Wolski, F., Ray, A., Schneider, J., Fong, R., Welinder, P., McGrew, B., Tobin, J., Abbeel, P., and Zaremba, W. (2017). Hindsight experience replay. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, pages 5055–5065.

[5] Baker, B., Gupta, O., Naik, N., and Raskar, R. (2016). Designing neural network architectures using reinforcement learning. In *ICLR (Poster)*.

[6] Barsce, J. C., Palombarini, J. A., and Martínez, E. C. (2017). Towards autonomous reinforcement learning: Automatic setting of hyper-parameters using bayesian optimization. In *2017 XLIII Latin American Computer Conference (CLEI)*, pages 1–9. IEEE.

[7] Baumann, M. and Buning, H. K. (2011). State aggregation by growing neural gas for reinforcement learning in continuous state spaces. In *2011 10th International Conference on Machine Learning and Applications and Workshops*, volume 1, pages 430–435. IEEE.

[8] Bello, I., Pham, H., Le, Q. V., Norouzi, M., and Bengio, S. (2016). Neural combinatorial optimization with reinforcement learning. *arXiv preprint arXiv:1611.09940*.

[9] Biedenkapp, A., Bozkurt, H. F., Eimer, T., Hutter, F., and Lindauer, M. (2020). Dynamic algorithm configuration: foundation of a new meta-algorithmic framework. In *Proceedings of the Twenty-fourth European Conference on Artificial Intelligence (ECAI'20)(Jun 2020)*.

[10] Bodnar, C., Day, B., and Lió, P. (2020). Proximal distilled evolutionary reinforcement learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 3283–3290.

[11] Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W. (2016a). Openai gym. *arXiv preprint arXiv:1606.01540*.

[12] Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W. (2016b). Openai gym.

[13] Chen, B., Zhang, K., Ou, L., Ba, C., Wang, H., and Wang, C. (2020). Automatic hyper-parameter optimization based on mapping discovery from data to hyper-parameters. *arXiv preprint arXiv:2003.01751*.

[14] Chiang, H.-T. L., Faust, A., Fiser, M., and Francis, A. (2019). Learning navigation behaviors end-to-end with autorl. *IEEE Robotics and Automation Letters*, 4(2):2007–2014.

[15] Chou, P.-W., Maturana, D., and Scherer, S. (2017). Improving stochastic policy gradients in continuous control with deep reinforcement learning using the beta distribution. In *International conference on machine learning*, pages 834–843. PMLR.

[16] Dabney, W., Ostrovski, G., Silver, D., and Munos, R. (2018). Implicit quantile networks for distributional reinforcement learning. In *International conference on machine learning*, pages 1096–1105. PMLR.

[17] Dai, H., Dai, B., and Song, L. (2016). Discriminative embeddings of latent variable models for structured data. In *International conference on machine learning*, pages 2702–2711.

[18] Daniel, C., Taylor, J., and Nowozin, S. (2016). Learning step size controllers for robust neural network training. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 30.

[19] Degroote, H., Bischl, B., Kotthoff, L., and De Causmaecker, P. (2016). Reinforcement learning for automatic online algorithm selection-an empirical study. *ITAT 2016 Proceedings*, 1649:93–101.

[20] Dong, X., Shen, J., Wang, W., Liu, Y., Shao, L., and Porikli, F. (2018). Hyperparameter optimization for tracking with continuous deep q-learning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 518–527.

[21] Elthakeb, A., Pilligundla, P., Mireshghallah, F., Yazdanbakhsh, A., Gao, S., and Esmaeilzadeh, H. (2019). Releq: An automatic reinforcement learning approach for deep quantization of neural networks. In *NeurIPS ML for Systems workshop, 2018*.

[22] Eysenbach, B., Gu, S., Ibarz, J., and Levine, S. (2018). Leave no trace: Learning to reset for safe and autonomous reinforcement learning. In *ICLR 2018 : International Conference on Learning Representations 2018*.

[23] Eysenbach, B., Salakhutdinov, R., and Levine, S. (2019). Search on the replay buffer: Bridging planning and reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 32, pages 15246–15257.

[24] Faust, A., Francis, A., and Mehta, D. (2019). Evolving rewards to automate reinforcement learning. *arXiv preprint arXiv:1905.07628*.

[25] Fernandez, F. C. and Caarls, W. (2018). Parameters tuning and optimization for reinforcement learning algorithms using evolutionary computing. In *2018 International Conference on Information Systems and Computer Science (INCISCOS)*, pages 301–305. IEEE.

[26] Finn, C., Abbeel, P., and Levine, S. (2017). Model-agnostic meta-learning for fast adaptation of deep networks. In *International Conference on Machine Learning*, pages 1126–1135. PMLR.

[27] Florensa, C., Held, D., Wulfmeier, M., Zhang, M., and Abbeel, P. (2017). Reverse curriculum generation for reinforcement learning. *Conference on Robot Learning*, pages 482–495.

[28] Forrester, A. I., Sóbester, A., and Keane, A. J. (2007). Multi-fidelity optimization via surrogate modelling. *Proceedings of the royal society a: mathematical, physical and engineering sciences*, 463(2088):3251–3269.

[29] Franke, J. K., Köhler, G., Biedenkapp, A., and Hutter, F. (2020). Sample-efficient automated deep reinforcement learning. *arXiv preprint arXiv:2009.01555*.

[30] Garcia, F. M. and Thomas, P. S. (2019). A meta-mdp approach to exploration for lifelong reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 32, pages 5691–5700.

[31] Grześ, M. (2017). Reward shaping in episodic reinforcement learning. In *AAMAS '17 Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems*, pages 565–573.

[32] Gu, S., Lillicrap, T., Sutskever, I., and Levine, S. (2016). Continuous deep q-learning with model-based acceleration. In *International Conference on Machine Learning*, pages 2829–2838. PMLR.

[33] Guo, X. and Hernández-Lerma, O. (2009). Continuous-time markov decision processes. In *Continuous-Time Markov Decision Processes*, pages 9–18. Springer.

[34] Haarnoja, T., Zhou, A., Hartikainen, K., Tucker, G., Ha, S., Tan, J., Kumar, V., Zhu, H., Gupta, A., Abbeel, P., et al. (2018). Soft actor-critic algorithms and applications. *arXiv preprint arXiv:1812.05905*.

[35] He, X., Zhao, K., and Chu, X. (2021). Automl: A survey of the state-of-the-art. *Knowledge-Based Systems*, 212:106622.

[36] Hutter, F., Hoos, H. H., and Leyton-Brown, K. (2011). Sequential model-based optimization for general algorithm configuration. In *International conference on learning and intelligent optimization*, pages 507–523. Springer.

[37] Ilavarasi, K. and Joseph, K. S. (2014). Variants of travelling salesman problem: A survey. In *International Conference on Information Communication and Embedded Systems (ICICES2014)*, pages 1–7. IEEE.

[38] Ivanovic, B., Harrison, J., Sharma, A., Chen, M., and Pavone, M. (2019). Barc: Backward reachability curriculum for robotic reinforcement learning. In *2019 International Conference on Robotics and Automation (ICRA)*, pages 15–21. IEEE.

[39] Jamieson, K. and Talwalkar, A. (2016). Non-stochastic best arm identification and hyperparameter optimization. In *Artificial Intelligence and Statistics*, pages 240–248. PMLR.

[40] Jomaa, H. S., Grabocka, J., and Schmidt-Thieme, L. (2019). Hyp-rl: Hyperparameter optimization by reinforcement learning. *arXiv preprint arXiv:1906.11527*.

[41] Judah, K., Fern, A., Tadepalli, P., and Goetschalckx, R. (2014). Imitation learning with demonstrations and shaping rewards. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 28.

[42] Kanervisto, A., Scheller, C., and Hautamäki, V. (2020). Action space shaping in deep reinforcement learning. In *2020 IEEE Conference on Games (CoG)*, pages 479–486. IEEE.

[43] Khadka, S. and Tumer, K. (2018). Evolution-guided policy gradient in reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 31, pages 1188–1200.

[44] Khalil, E., Dai, H., Zhang, Y., Dilkina, B., and Song, L. (2017). Learning combinatorial optimization algorithms over graphs. In *Advances in Neural Information Processing Systems*, pages 6348–6358.

[45] Laroche, R. and Feraud, R. (2018). Reinforcement learning algorithm selection. In *ICLR 2018 : International Conference on Learning Representations 2018*.

[46] Li, L., Jamieson, K., DeSalvo, G., Rostamizadeh, A., and Talwalkar, A. (2017). Hyperband: A novel bandit-based approach to hyperparameter optimization. *The Journal of Machine Learning Research*, 18(1):6765–6816.

[47] Lillicrap, T. P., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., Silver, D., and Wierstra, D. (2016). Continuous control with deep reinforcement learning. In *ICLR 2016 : International Conference on Learning Representations 2016*.

[48] Lin, S. and Wright, R. (2010). Evolutionary tile coding: An automated state abstraction algorithm for reinforcement learning. In *Proceedings of the 8th AAAI Conference on Abstraction, Reformulation, and Approximation*, pages 42–47.

[49] Maclaurin, D., Duvenaud, D., and Adams, R. (2015). Gradient-based hyperparameter optimization through reversible learning. In *International conference on machine learning*, pages 2113–2122. PMLR.

[50] Mazyavkina, N., Sviridov, S., Ivanov, S., and Burnaev, E. (2020). Reinforcement learning for combinatorial optimization: A survey. *arXiv preprint arXiv:2003.03600*.

[51] Miikkulainen, R., Liang, J., Meyerson, E., Rawal, A., Fink, D., Francon, O., Raju, B., Shahrzad, H., Navruzyan, A., Duffy, N., et al. (2019). Evolving deep neural networks. In *Artificial intelligence in the age of neural networks and brain computing*, pages 293–312. Elsevier.

[52] Mnih, V., Badia, A. P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., Silver, D., and Kavukcuoglu, K. (2016). Asynchronous methods for deep reinforcement learning. In *International conference on machine learning*, pages 1928–1937. PMLR.

[53] Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., et al.

(2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533.

[54] Ng, A. Y., Harada, D., and Russell, S. (1999). Policy invariance under reward transformations: Theory and application to reward shaping. In *Icml*, volume 99, pages 278–287.

[55] Ottoni, A. L., Nepomuceno, E. G., de Oliveira, M. S., and de Oliveira, D. C. (2020). Tuning of reinforcement learning parameters applied to sop using the scott–knott method. *Soft Computing*, 24(6):4441–4453.

[56] Peer, E., Menkovski, V., Zhang, Y., and Lee, W.-J. (2018). Shunting trains with deep reinforcement learning. In *2018 ieee international conference on systems, man, and cybernetics (smc)*, pages 3063–3068. IEEE.

[57] Perozzi, B., Al-Rfou, R., and Skiena, S. (2014). Deepwalk: Online learning of social representations. In *Proceedings of the 20th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 701–710.

[58] Peters, J., Mulling, K., and Altun, Y. (2010). Relative entropy policy search. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 24.

[59] Puterman, M. L. (2014). *Markov decision processes: discrete stochastic dynamic programming.* John Wiley & Sons.

[60] Raschka, S. (2018). Model evaluation, model selection, and algorithm selection in machine learning. *arXiv preprint arXiv:1811.12808*.

[61] Refaei Afshar, R., Rhuggenaath, J., Zhang, Y., and Kaymak, U. (2021). A reward shaping approach for reserve price optimization using deep reinforcement learning. In *The International Joint Conference on Neural Networks (IJCNN2021)*.

[62] Refaei Afshar, R., Zhang, Y., Firat, M., and Kaymak, U. (2019). A reinforcement learning method to select ad networks in waterfall strategy. In

11th International Conference on Agents and Artificial Intelligence, ICAART 2019, pages 256–265. SCITEPRESS-Science and Technology Publications, Lda.

[63] Refaei Afshar, R., Zhang, Y., Firat, M., and Kaymak, U. (2020). A state aggregation approach for solving knapsack problem with deep reinforcement learning. *arXiv preprint arXiv:2004.12117*.

[64] Scarselli, F., Gori, M., Tsoi, A. C., Hagenbuchner, M., and Monfardini, G. (2008). The graph neural network model. *IEEE transactions on neural networks*, 20(1):61–80.

[65] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.

[66] Sehgal, A., La, H., Louis, S., and Nguyen, H. (2019). Deep reinforcement learning using genetic algorithm for parameter optimization. In *2019 Third IEEE International Conference on Robotic Computing (IRC)*, pages 596–601. IEEE.

[67] Sherstov, A. A. and Stone, P. (2005). Function approximation via tile coding: Automating parameter choice. In *International Symposium on Abstraction, Reformulation, and Approximation*, pages 194–205. Springer.

[68] Singh, S., Barto, A. G., and Chentanez, N. (2004). Intrinsically motivated reinforcement learning. In *Proceedings of the 17th International Conference on Neural Information Processing Systems*, pages 1281–1288.

[69] Smart, W. D. and Kaelbling, L. P. (2002). Effective reinforcement learning for mobile robots. In *Proceedings 2002 IEEE International Conference on Robotics and Automation (Cat. No. 02CH37292)*, volume 4, pages 3404–3410. IEEE.

[70] Spaan, M. T. (2012). Partially observable markov decision processes. In *Reinforcement Learning*, pages 387–414. Springer.

[71] Stanley, K. O., Clune, J., Lehman, J., and Miikkulainen, R. (2019). Designing neural networks through neuroevolution. *Nature Machine Intelligence*, 1(1):24–35.

[72] Su, P.-H., Vandyke, D., Gasic, M., Mrksic, N., Wen, T.-H., and Young, S. (2015). Reward shaping with recurrent neural networks for speeding up on-line policy learning in spoken dialogue systems. In *Proceedings of the 16th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 417–421.

[73] Sutton, R. S. (1996). Generalization in reinforcement learning: Successful examples using sparse coarse coding. *Advances in neural information processing systems*, pages 1038–1044.

[74] Sutton, R. S. and Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

[75] Tang, Y. and Agrawal, S. (2020). Discretizing continuous action space for on-policy optimization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 5981–5988.

[76] Tavakoli, A., Fatemi, M., and Kormushev, P. (2020). Learning to represent action values as a hypergraph on the action vertices. In *International Conference on Learning Representations*.

[77] Tessler, C., Tennenholtz, G., and Mannor, S. (2019). Distributional policy optimization: An alternative approach for continuous control. *Advances in Neural Information Processing Systems*, 32:1352–1362.

[78] Vinyals, O., Fortunato, M., and Jaitly, N. (2015). Pointer networks. In *Proceedings of the 28th International Conference on Neural Information Processing Systems-Volume 2*, pages 2692–2700.

[79] Vuong, Q., Vikram, S., Su, H., Gao, S., and Christensen, H. I. (2019). How to pick the domain randomization parameters for sim-to-real transfer of reinforcement learning policies? *arXiv preprint arXiv:1903.11774*.

[80] Wang, J. X., Kurth-Nelson, Z., Tirumala, D., Soyer, H., Leibo, J. Z., Munos, R., Blundell, C., Kumaran, D., and Botvinick, M. (2016). Learning to reinforcement learn. *Cognitive Science.*

[81] Wang, Z., Bapst, V., Heess, N., Mnih, V., Munos, R., Kavukcuoglu, K., and de Freitas, N. (2016). Sample efficient actor-critic with experience replay. *arXiv preprint arXiv:1611.01224.*

[82] White, M. and White, A. (2016). A greedy approach to adapting the trace parameter for temporal difference learning. In *Proceedings of the 2016 International Conference on Autonomous Agents & Multiagent Systems*, pages 557–565.

[83] Whiteson, S. (2010). Adaptive tile coding. In *Adaptive Representations for Reinforcement Learning*, pages 65–76. Springer.

[84] Wiering, M. and Van Otterlo, M. (2012). Reinforcement learning. *Adaptation, learning, and optimization*, 12:3.

[85] Xu, C., Qin, T., Wang, G., and Liu, T.-Y. (2017). Reinforcement learning for learning rate control. *arXiv preprint arXiv:1705.11159.*

[86] Zhou, C., Liu, Y., Liu, X., Liu, Z., and Gao, J. (2017). Scalable graph embedding for asymmetric proximity. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 31.

[87] Zoph, B. and Le, Q. V. (2016). Neural architecture search with reinforcement learning. In *ICLR.*

[88] Zoph, B., Vasudevan, V., Shlens, J., and Le, Q. V. (2018). Learning transferable architectures for scalable image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 8697–8710.

[89] Zou, H., Ren, T., Yan, D., Su, H., and Zhu, J. (2019). Reward shaping via meta-learning. *arXiv preprint arXiv:1901.09330.*